(a) Today: one flagship call over a fixed 30-candidate shortlist

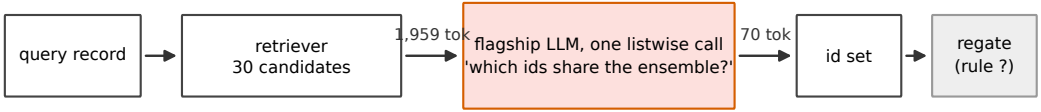

$0.020 / query (GPT-4.1) · fixed K · no abstention · no verification · title still shown at L5

(b) Proposed v2 system: fuse what the retriever knows, read less, say 'I do not know', verify under permutation

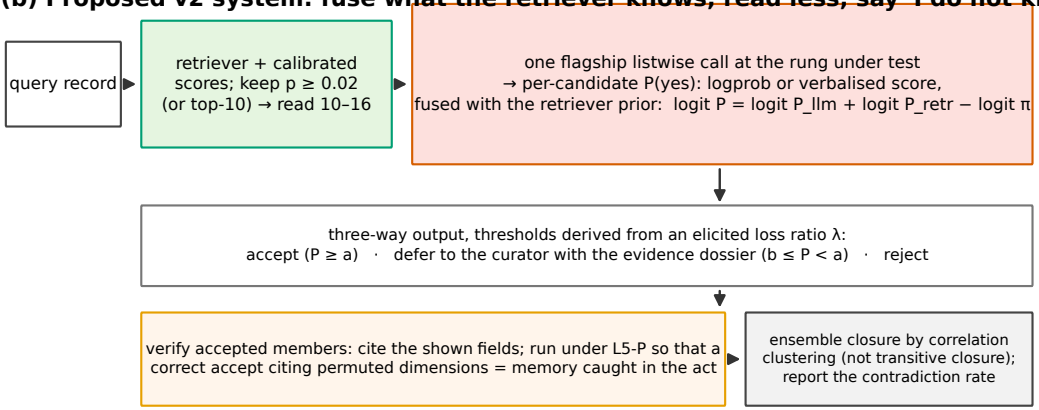

cost ≈ 0.94–1.24 ¢/query (flagship, K=10 or prune, + prior) · ≈ 0.19 ¢ (mini variant) · vs 1.99 ¢ today — real token prices; quality: simulation only

gates before adopting (b): the $0 fusion test on logged scores and the $1 K=10-vs-K=30 test; if both are null, keep (a) and spend on the measurement

the ladder MEASUREMENT uses the fixed-rung system (a) at every rung; system (b) is evaluated end-to-end on its own quality–cost point, never pooled into it