(a) Today: one flagship call over a fixed 30-candidate shortlist

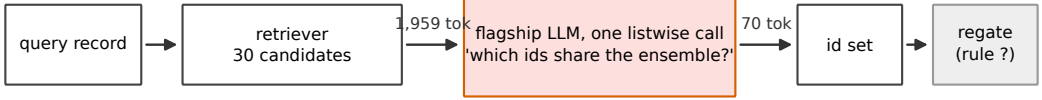

$0.020 / query (GPT-4.1) · fixed K · no abstention · no verification · title still shown at L5

(b) Proposed v2: read less, escalate only what is open, test every claim

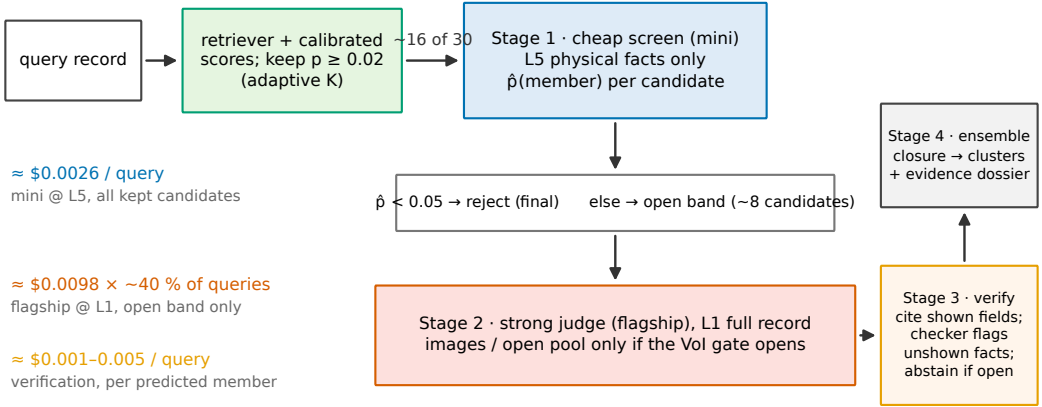

all-in ≈ $0.0075–0.0115 / query (−40 … −63 % vs today) · share of decisions made from physical facts = a per-decision memorisation metric