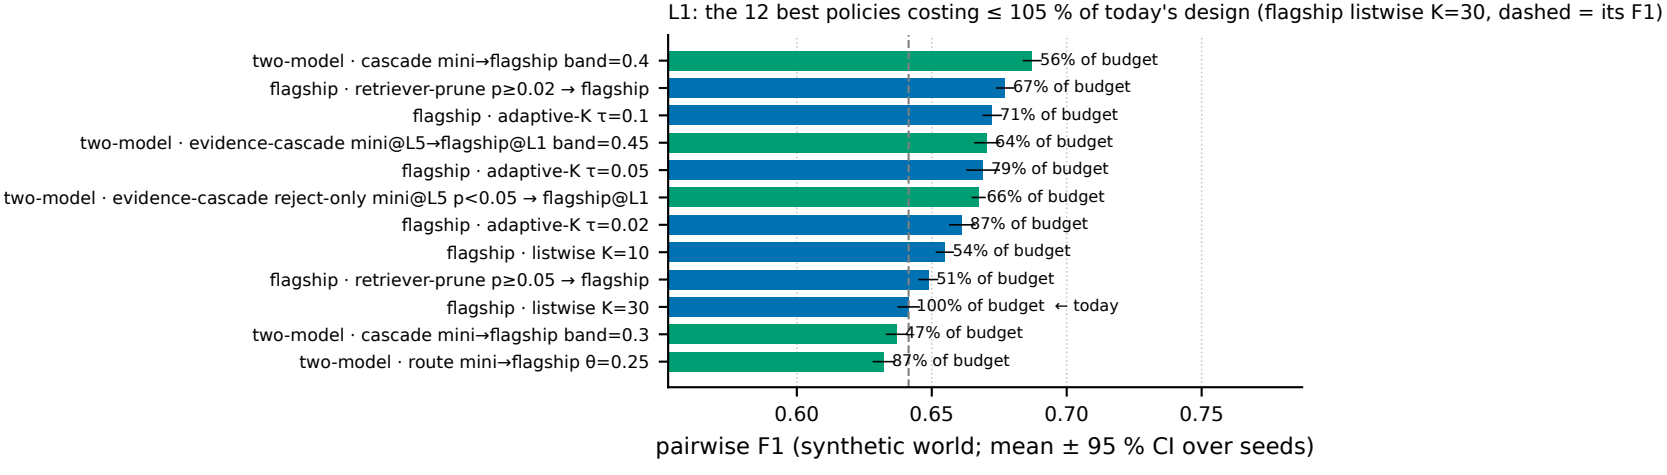
L1: the 12 best policies costing ≤ 105 % of today's design (flagship listwise K=30, dashed = its F1)
two-model · cascade mini→flagship band=0.4
flagship · retriever-prune p≥0.02 → flagship
flagship · adaptive-K τ=0.1
two-model · evidence-cascade mini@L5→flagship@L1 band=0.45
flagship · adaptive-K τ=0.05
two-model · evidence-cascade reject-only mini@L5 p<0.05 → flagship@L1
flagship · adaptive-K τ=0.02
flagship · listwise K=10
flagship · retriever-prune p≥0.05 → flagship
flagship · listwise K=30
two-model · cascade mini→flagship band=0.3
two-model · route mini→flagship θ=0.25
56% of budget
67% of budget
71% of budget
64% of budget
79% of budget
66% of budget
87% of budget
54% of budget
51% of budget
100% of budget ← today
47% of budget
87% of budget
0.60
0.65
0.70
0.75
pairwise F1 (synthetic world; mean ± 95 % CI over seeds)

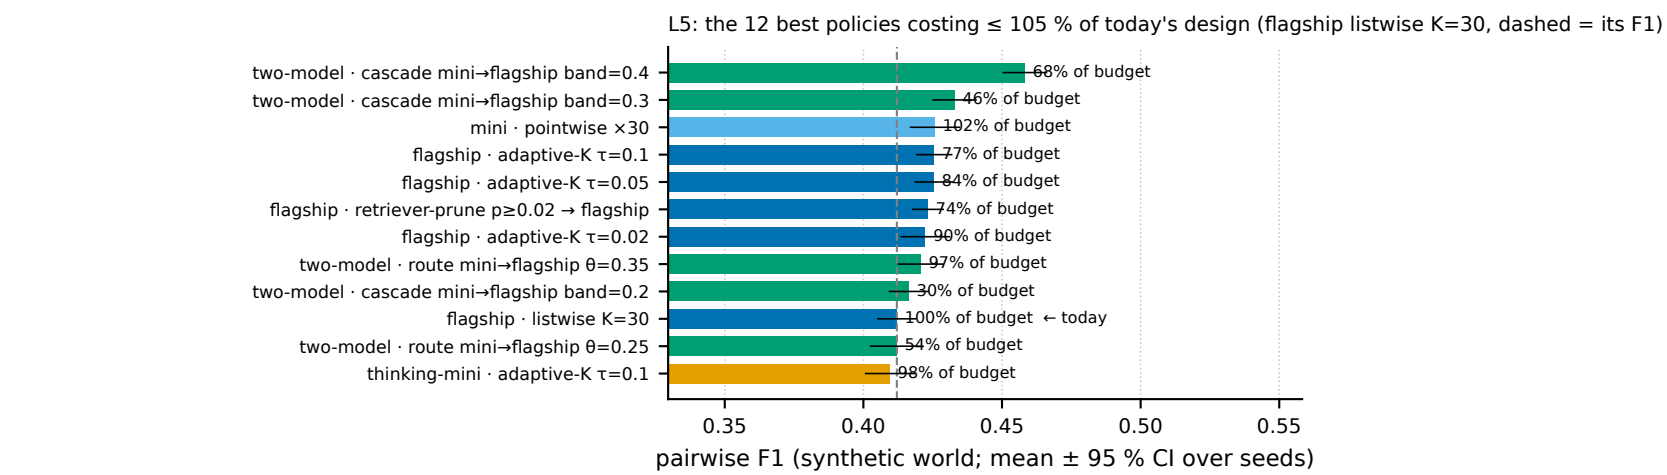
L5: the 12 best policies costing ≤ 105 % of today's design (flagship listwise K=30, dashed = its F1)
two-model · cascade mini→flagship band=0.4
two-model · cascade mini→flagship band=0.3
mini · pointwise ×30
flagship · adaptive-K τ=0.1
flagship · adaptive-K τ=0.05
flagship · retriever-prune p≥0.02 → flagship
flagship · adaptive-K τ=0.02
two-model · route mini→flagship θ=0.35
two-model · cascade mini→flagship band=0.2
flagship · listwise K=30
two-model · route mini→flagship θ=0.25
thinking-mini · adaptive-K τ=0.1
68% of budget
46% of budget
102% of budget
77% of budget
84% of budget
74% of budget
90% of budget
97% of budget
30% of budget
100% of budget ← today
54% of budget
98% of budget
0.35
0.40
0.45
0.50
0.55
pairwise F1 (synthetic world; mean ± 95 % CI over seeds)