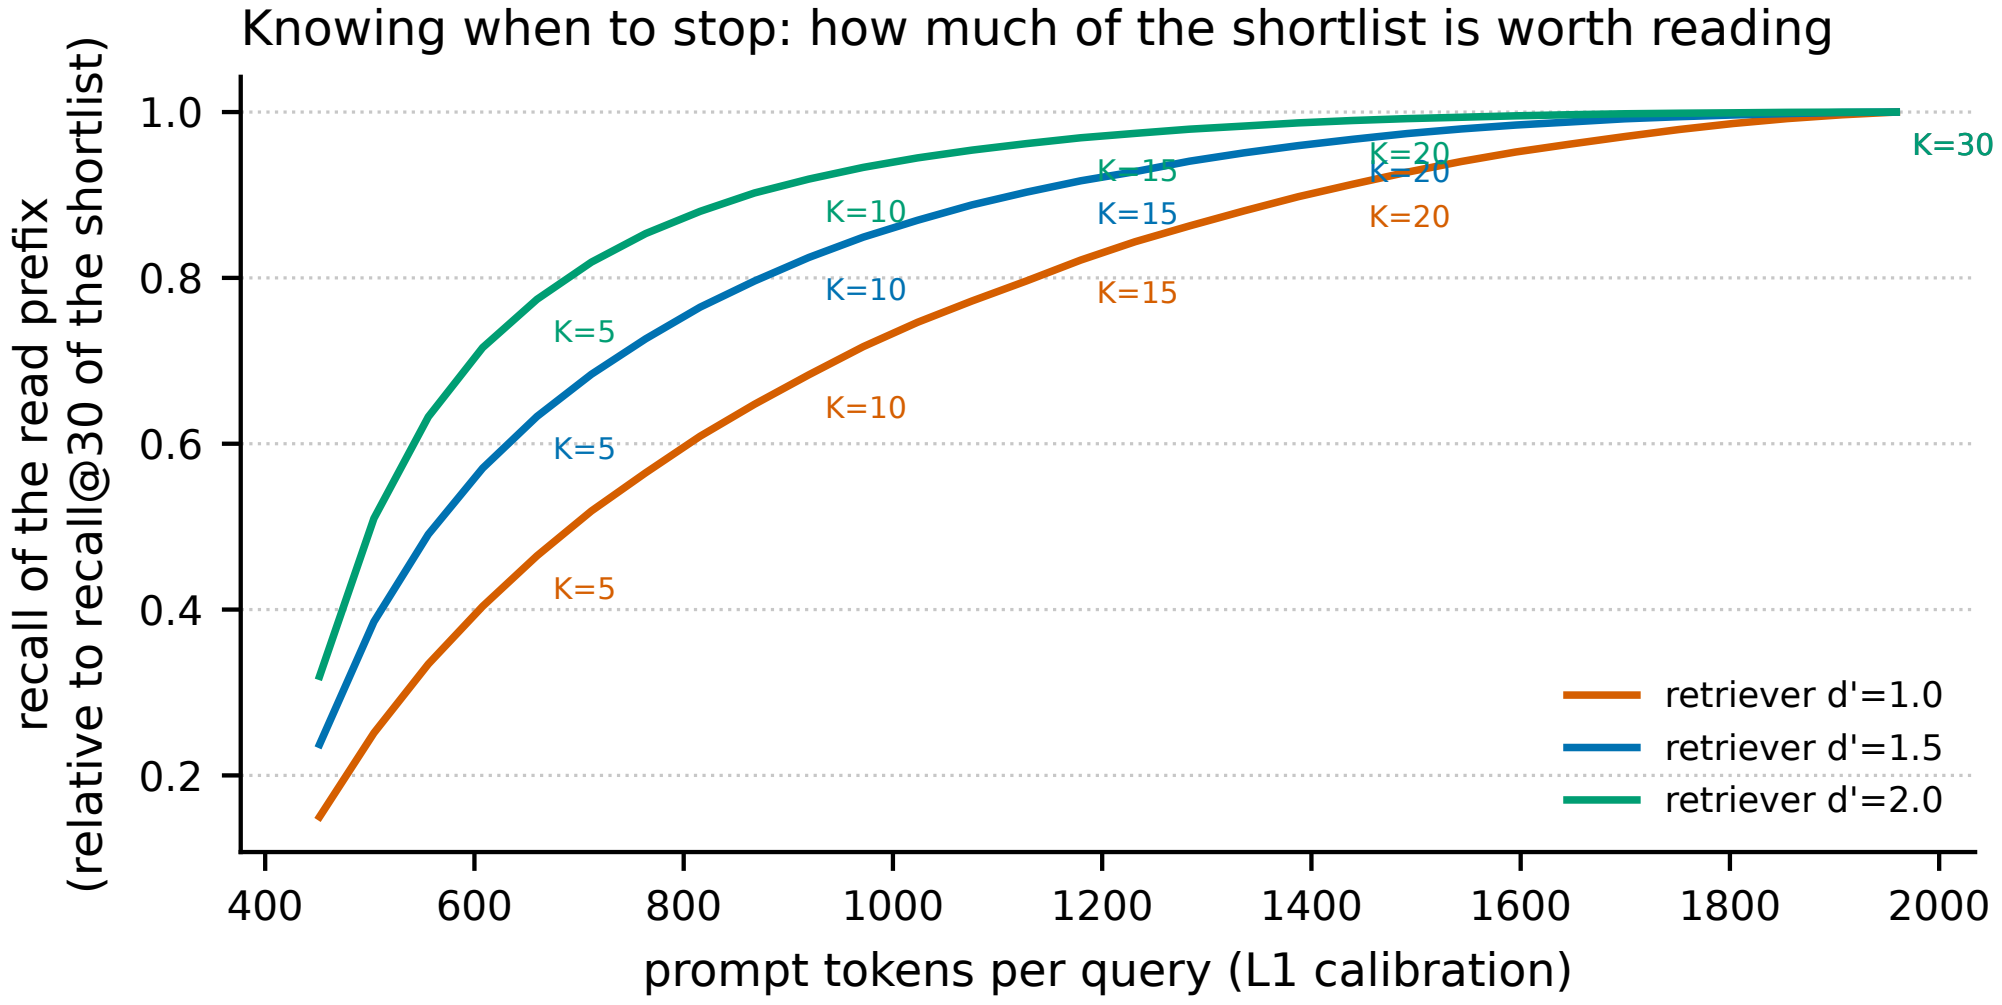
Knowing when to stop: how much of the shortlist is worth reading
recall of the read prefix
(relative to recall@30 of the shortlist)
1.0
0.8
0.6
0.4
0.2
400
600
800
1000
1200
1400
1600
1800
2000
prompt tokens per query (L1 calibration)
K=5
K=10
K=15
K=20
K=30
K=5
K=10
K=15
K=20
K=5
K=10
K=15
K=20
retriever d'=1.0
retriever d'=1.5
retriever d'=2.0